

MLT Homework 8

Ana Borovac
Bas Haver

November 12, 2018

Question 1

Let $\psi(\lambda) = \frac{\lambda^2}{2}$. The Legendre-Fenchel transform of ψ is given by

$$\psi^*(\epsilon) = \sup_{\lambda \in \mathbb{R}} \lambda\epsilon - \psi(\lambda).$$

Subquestion 1.1

$$\psi^*(\epsilon) = \frac{\epsilon^2}{2}.$$

Solution

Let define $\psi_1(\lambda)$:

$$\psi_1(\lambda) = \lambda\epsilon - \frac{\lambda^2}{2}$$

Furthermore:

$$\psi'_1(\lambda) = \epsilon - \lambda$$

Since $\psi_1(\lambda)$ is a parabola it has only one extreme; particularly it has just a maximum (negative sign before λ^2). So, the maximum is reached at:

$$\lambda = \epsilon \quad \Rightarrow \quad \psi_1(\epsilon) = \epsilon \cdot \epsilon - \frac{\epsilon^2}{2} = \frac{\epsilon^2}{2}$$

We can conclude:

$$\psi^*(\epsilon) = \psi_1(\epsilon) = \frac{\epsilon^2}{2}$$

Subquestion 1.2

$$(\psi^*)^{-1}(z) = \pm\sqrt{2z}.$$

Solution

From previous point we know:

$$\psi^*(\epsilon) = \frac{\epsilon^2}{2}$$

It follows:

$$\begin{aligned} z &= \frac{\epsilon^2}{2} \\ 2z &= \epsilon^2 \\ \epsilon &= \pm\sqrt{2z} \end{aligned}$$

So:

$$(\psi^*)^{-1}(z) = \pm\sqrt{2z}$$

Question 2

The Blooper Reel

Subquestion 2.1

Deterministic fails for Adversarial Bandits Show that any deterministic algorithm (UCB included) has linear regret in the adversarial bandit setting. Hint: you can use the argument on the top of page 23.

Solution

We were a bit confused by this question, since it states that we need to find a linear regret, which we did not find. Instead we found that it can be bounded between two linear functions, but it does not necessarily is linear itself.

As a counter-example of the linearity we have that for $n = 0$ no regret has been obtained. Therefore linearity only holds when $R_m + R_n = R_{n+m}$ for all n, m . But when we choose $n = 4$, four arms and choose to play on arms 1,2,3 and 4 succesively, we find $R_1 = 1$, $R_3 = 3$, $R_4 = 3$, so linearity does not hold. It does however hold that it remains between our bounds $\frac{n}{K}$ and n for K the amount of arms.

We use the hint on the top of page 23, which tells us that for a deterministic forecaster, we can use the following sequence of losses:

$$\begin{aligned} \text{if } I_t = 1, \quad & \text{then } l_{2,t} = 0 \quad \text{and} \quad l_{i,t} = 1 \quad \text{for all } i \neq 2; \\ \text{if } I_t \neq 1, \quad & \text{then } l_{1,t} = 0 \quad \text{and} \quad l_{i,t} = 1 \quad \text{for all } i \neq 1 \end{aligned}$$

Of course this is just a worst-case feedback. For every choice of arm, we will get a loss of 1, which result in a regret that is as high as possible. Since it does the

trick, we will use it.

This sequence of losses now implies the following for the regret:

$$\begin{aligned} R_n &= \sum_{t=1}^n l_{t, I_t} - \min_k \sum_{t=1}^n l_n^k \\ &= n - \min_k \sum_{t=1}^n l_n^k \end{aligned}$$

But now for any choice of arm, with at least two arms, $\min_k \sum_{t=1}^n l_n^k$ is at most $\frac{n}{K}$. Therefore the regret is bounded from below by $\frac{n}{2}$ and since the loss function is nonnegative, the regret is also bounded from above by n .

Subquestion 2.2

Consider a K -armed stochastic bandit model with unit-variance Gaussian rewards with means μ_1, \dots, μ_K . In round t the learner chooses arm $I_t \in [K]$ and receives reward $X_t \sim \mathcal{N}(\mu_{I_t}, 1)$, where μ_i is the (unknown) reward of arm i . Now let's fix the following algorithm, which is inspired by Empirical Risk Minimisation:

- (a) First, pull every arm once (that is $T_i = t$ for $t \leq K$).
- (b) Then after each number $t \geq K$ of rounds, from the empirical estimates

$$\hat{\mu}_i(t) = \frac{\sum_{s=1}^t \mathbb{1}_{\{I_s=i\}} X_s}{\sum_{s=1}^t \mathbb{1}_{\{I_s=i\}}}$$

and play $I_{t+1} = \arg \max_i \hat{\mu}_i(t)$.

For $K = 2$, show that this algorithm has pseudo-regret

$$\bar{R} = n\mu^* - \mathbb{E}[\sum_{t=1}^n \mu_{I_t}]$$

that is *linear* in n .

Hint: you can use the following outline. Assume $\mu_1 > \mu_2$. Pick some threshold $\epsilon > 0$ (which you will optimise in a later step).

- Argue that with constant probability (independent of n) the reward drawn from the best arm in the first phase is below $\mu_2 - \epsilon$.
- Bound the probability that for a single time step t we have $\hat{\mu}_2(t) < \mu_2 - \epsilon$ using Chernoff's bound.
- Use the union bound to bound the probability that $\exists t \geq 2 : \hat{\mu}_2(t) < \mu_2 - \epsilon$.
- Now pick ϵ large enough so that the previous probability bound is non-trivial (i.e. is ≥ 1).

Conclude that with some small probability the sample from the best arm is very low, and the samples from the second-best arm are all typical, so the algorithm keeps pulling arm 2 only. Deduce that the pseudo-regret is hence linear in n .

Solution

We will follow the outline as in the hint. So let's assume $\mu_1 > \mu_2$ and let $\epsilon > 0$ to be chosen later. The first phase is determined by pulling every arm once, so in our case pull the first arm and then pull the second arm. Now μ_1 and μ_2 are estimated by setting the estimator equal to the value we have found. The result of this part is not dependent of whatever happens after this, so certainly not the amount of pulls. Thus the probability that the reward from the best arm in the first phase is below $\mu_2 - \epsilon$ is independent of n , let us call this constant in n to be α , which is positive since we observe a normal distribution. Next we will find an upper bound for

$$\mathbb{P}(\hat{\mu}_2(t) < \mu_2 - \epsilon \mid T_2(t) = s),$$

where T_2 is the amount of times that the second arm was pulled. For this we rewrite $\{\hat{\mu}_2(t) < \mu_2 - \epsilon\} = \{\hat{\mu}_2(t) - \mu_2 < -\epsilon\}$ to be able to use Chernoff's bound:

$$\begin{aligned} \mathbb{P}(\hat{\mu}_2(t) < \mu_2 - \epsilon \mid T_2(t) = s) &= \mathbb{P}(\hat{\mu}_2(t) - \mu_2 < -\epsilon \mid T_2(t) = s) \\ &\leq e^{-s \frac{\epsilon^2}{2}} \end{aligned}$$

which holds for all t and s .

Now we find the probability that there is a $t \geq 2$ for which $\hat{\mu}_2(t) < \mu_2 - \epsilon$ to be smaller or equal than the probability of finding this after the first phase plus the probability of not finding it after this phase but finding it later. This latter happens with a probability which is bounded as we have seen above by $e^{-s \frac{\epsilon^2}{2}}$. This gives

$$\mathbb{P}(\exists t \geq 2 : \hat{\mu}_2(t) < \mu_2 - \epsilon) \leq \alpha + e^{-s \frac{\epsilon^2}{2}}$$

Now we will choose our ϵ appropriately to obtain a non-trivial probability bound. Therefore we notice that s is at least 1, since we pulled the second machine at the first phase. This gives now for $\epsilon > \sqrt{-2 \log(1 - \alpha)}$:

$$\begin{aligned} \alpha + e^{-s \frac{\epsilon^2}{2}} &\leq \alpha + e^{-\frac{\epsilon^2}{2}} \\ &< \alpha + e^{-\frac{\sqrt{-2 \log(1 - \alpha)}^2}{2}} \\ &= \alpha + e^{-\frac{-2 \log(1 - \alpha)}{2}} \\ &= \alpha + e^{\log(1 - \alpha)} \\ &= 1 \end{aligned}$$

Now by the central limit theorem, the probability of finding a very low sample of the best arm is very low on the long run and since the samples of the second-best arm are all typical, the algorithm keeps pulling arm 2 only. Then we also have $\mathbb{E}[\sum_{t=1}^n \mu_{I_t}] = \mathbb{E}[\sum_{t=1}^n \mu_2]$ in behaviour, which gives us that the pseudo-regret is linear in n .

Question 3

We consider an adversarial bandit model with K^2 arms indexed by $i \in [K]$ and $j \in [K]$. For each arm (i, j) , the loss at time t is $a_t^i + b_t^j$, where $a_t^i \in [0, 1]$ and $b_t^j \in [0, 1]$ are chosen by the adversary before the start of the interaction. Then each round the learner picks an arm $(I_t, J_t) \in [K]^2$ and observes $a_t^{I_t}$ and $b_t^{J_t}$ separately (and incurs their sum as the loss).

Subquestion 3.1

Consider running a single instance of EXP3 on all K^2 arms (with loss range $[0, 2]$). Show that the expected pseudo-regret compared to the best arm (i^*, j^*) is bounded by

$$\bar{R}_n \leq 2\sqrt{2nK^2 \ln(K^2)}$$

Solution

Below we used the following facts:

- $\min x + y = \min x + \min y$; $x, y \geq 0$
- Linearity of expected value.
- Theorem from the lectures: $\bar{R}_n \leq \sqrt{2nK \ln K}$, where K is the number of arms; in our case we have K^2 arms.

$$\begin{aligned} \bar{R}_n &= \mathbb{E}_{I_1, \dots, I_n, J_1, \dots, J_n} \left\{ \sum_{t=1}^n a_t^{I_t} + b_t^{J_t} \right\} - \min_k \sum_{t=1}^n a_t^k + b_t^k \\ &= \left(\mathbb{E}_{I_1, \dots, I_n} \left\{ \sum_{t=1}^n a_t^{I_t} \right\} - \min_k \sum_{t=1}^n a_t^k \right) + \left(\mathbb{E}_{J_1, \dots, J_n} \left\{ \sum_{t=1}^n b_t^{J_t} \right\} - \min_k \sum_{t=1}^n b_t^k \right) \\ &\leq \sqrt{2nK^2 \ln K^2} + \sqrt{2nK^2 \ln K^2} \\ &= 2\sqrt{2nK^2 \ln K^2} \end{aligned}$$

Subquestion 3.2

Now we will use the a_t^i and b_t^j observations separately. Consider running two K -arm instances of EXP3, one with $i \rightarrow a_t^i$ as the loss and one with $j \rightarrow b_t^j$ as the loss. Have the first algorithm control I_t and the second J_t . Show that the overall expected pseudo-regret is bounded by

$$\bar{R}_n \leq 2\sqrt{2nK \ln K}.$$

Solution

We do similar as before, just that now we choose arm in the set of K arms and not K^2 .

$$\begin{aligned}\bar{R}_n &= \mathbb{E}_{I_1, \dots, I_n, J_1, \dots, J_n} \left\{ \sum_{t=1}^n a_t^{I_t} + b_t^{J_t} \right\} - \min_k \sum_{t=1}^n a_t^k + b_t^k \\ &= \left(\mathbb{E}_{I_1, \dots, I_n} \left\{ \sum_{t=1}^n a_t^{I_t} \right\} - \min_k \sum_{t=1}^n a_t^k \right) + \left(\mathbb{E}_{J_1, \dots, J_n} \left\{ \sum_{t=1}^n b_t^{J_t} \right\} - \min_k \sum_{t=1}^n b_t^k \right) \\ &\leq \sqrt{2nK \ln K} + \sqrt{2nK \ln K} \\ &= 2\sqrt{2nK \ln K}\end{aligned}$$