

MLT Homework 10

Ana Borovac
Bas Haver

November 19, 2018

Question 1

Revisiting the AA

Subquestion 1.1

Non-uniform regret bound for AA with prior Fix a prior distribution $\pi \in \Delta_K$. Consider the Aggregating Algorithm with prior π , defined by:

$$w_t^k = \frac{\pi^k e^{-\sum_{s=1}^{t-1} l_s^k}}{\sum_{j=1}^K \pi^j e^{-\sum_{s=1}^{t-1} l_s^j}} \quad (1)$$

Show that for the mix-loss game, the regret w.r.t. each expert k is at most

$$R_T^k = \sum_{t=1}^T (\hat{l}_t - l_t^k) \leq -\ln \pi^k$$

Solution

In the final solution we needed the following fact:

$$\begin{aligned} -\ln \left(\sum_{k=1}^K \pi^k e^{-\sum_{s=1}^T l_s^k} \right) &\leq \sum_{s=1}^T l_s^k - \ln \pi^k \\ \ln \left(\sum_{k=1}^K \pi^k e^{-\sum_{s=1}^T l_s^k} \right) &\geq -\sum_{s=1}^T l_s^k + \ln \pi^k \\ \sum_{k=1}^K \pi^k e^{-\sum_{s=1}^T l_s^k} &\geq e^{-\sum_{s=1}^T l_s^k + \ln \pi^k} \\ \sum_{k=1}^K \pi^k e^{-\sum_{s=1}^T l_s^k} &\geq \pi^k e^{-\sum_{s=1}^T l_s^k}; \quad \forall k \end{aligned}$$

We also used:

$$\sum_{j=1}^K \pi^j = 1 \quad \Rightarrow \quad \ln \left(\sum_{j=1}^K \pi^j \right) = 0$$

So:

$$\begin{aligned} \sum_{t=1}^T \hat{l}_t &= \sum_{t=1}^T -\ln \left(\sum_{k=1}^K w_t^k e^{-l_t^k} \right) \\ &= \sum_{t=1}^T -\ln \left(\sum_{k=1}^K \frac{\pi^k e^{-\sum_{s=1}^{t-1} l_s^k}}{\sum_{j=1}^K \pi^j e^{-\sum_{s=1}^{t-1} l_s^j}} e^{-l_t^k} \right) \\ &= \sum_{t=1}^T -\ln \left(\frac{\sum_{k=1}^K \pi^k e^{-\sum_{s=1}^t l_s^k}}{\sum_{j=1}^K \pi^j e^{-\sum_{s=1}^{t-1} l_s^j}} \right) \\ &= -\ln \prod_{t=1}^T \left(\frac{\sum_{k=1}^K \pi^k e^{-\sum_{s=1}^t l_s^k}}{\sum_{j=1}^K \pi^j e^{-\sum_{s=1}^{t-1} l_s^j}} \right) \\ &= -\ln \left(\frac{\sum_{k=1}^K \pi^k e^{-\sum_{s=1}^T l_s^k}}{\sum_{j=1}^K \pi^j} \right) \\ &= -\ln \left(\sum_{k=1}^K \pi^k e^{-\sum_{s=1}^T l_s^k} \right) + \ln \left(\sum_{j=1}^K \pi^j \right) \\ &\leq \sum_{s=1}^T l_s^k - \ln \pi^k \\ \Rightarrow \quad \sum_{t=1}^T \hat{l}_t - l_t^k &\leq -\ln \pi^k \end{aligned}$$

Subquestion 1.2

The AA does not exploit non-stationarity Consider the Aggregating Algorithm in its incremental representation:

$$w_1^k = \frac{1}{K} \quad \text{and} \quad w_{t+1}^k = w_t^k e^{-l_t^k}.$$

At first sight, the AA adapts its weights sequentially to better fit the instantaneous losses. So maybe it can already compete with sequences of experts! Show that this is not the case, by constructing a sequence of losses such that the regret of the AA compared to the best sequence with one switch is arbitrarily high, even at a fixed horizon T .

Solution

We are observing sequences of form:

$$\xi \in [K]^T : \xi = (\underbrace{k, \dots, k}_m, \underbrace{j, \dots, j}_n); \quad \text{for } k \neq j$$

- If

Question 2

Collapsed Fixed Share Computation

In the Notes, we defined Fixed Share to predict with

$$w_t^k := \sum_{\xi \in [K]^T : \xi_t = k} w_t^\xi = \sum_{\xi \in [K]^T : \xi_t = k} \frac{\pi^\xi e^{-\sum_{s=1}^{t-1} l_s^\xi}}{\sum_{j \in [K]^T} \pi^j e^{-\sum_{s=1}^{t-1} l_s^j}}$$

where w_t^ξ are the weights of the AA (1) when run on experts $[K]^T$ with prior

$$\pi^\xi := \frac{1}{K} \prod_{t=2}^T \begin{cases} 1 - \alpha & \text{if } \xi_{t-1} = \xi_t \\ \frac{\alpha}{K-1} & \text{if } \xi_{t-1} \neq \xi_t \end{cases}$$

and losses defined by $l_t^\xi = l_t^{\xi^k}$. In this exercise we find a way to maintain w_t^k directly in $O(K)$ time per round.

Subquestion 2.1

What is w_1^k ?

Solution

$$\begin{aligned} w_1^k &= \sum_{\xi \in [K]^T : \xi_1 = k} w_1^\xi \\ &= \sum_{\xi \in [K]^T : \xi_1 = k} \pi^\xi \\ &= \sum_{\xi \in [K]^T : \xi_1 = k} \frac{1}{K} \\ &= |\xi \in [K]^T : \xi_1 = k| \frac{1}{K} \end{aligned}$$

Subquestion 2.2

Show that the Fixed Share weights satisfy the recurrence

$$w_{t+1}^k = (1 - \beta) \frac{w_t^k e^{-l_t^k}}{\sum_j w_t^j e^{-l_t^j}} + \frac{\beta}{K}$$

where $\beta = \alpha \frac{K}{K-1}$.

Solution

First, let us prove the hint. So, we would like to see that the following holds:

$$\sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi = \begin{cases} (1 - \alpha) \sum_{\xi: \xi_t=j} w_t^\xi, & j = k \\ \frac{\alpha}{K-1} \sum_{\xi: \xi_t=j} w_t^\xi, & j \neq k \end{cases}$$

For any $k \neq j$ it holds:

$$\begin{aligned} \frac{\sum_{\xi: \xi_t=j, \xi_{t+1}=j} w_t^\xi}{\sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi} &= \frac{\sum_{\xi: \xi_t=j, \xi_{t+1}=j} \frac{1}{K} \prod^* (1 - \alpha) \prod^* e^{-\Sigma^*}}{\sum_{\xi: \xi_t=j, \xi_{t+1}=j} \frac{1}{K} \prod^* \frac{\alpha}{K-1} \prod^* e^{-\Sigma^*}} \\ &= \frac{1 - \alpha}{\frac{\alpha}{K-1}} \end{aligned}$$

We also know:

$$\sum_{k=1}^K \sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi = \sum_{\xi_t=j} w_t^\xi$$

From that it follows ($k \neq j$):

$$\begin{aligned} \sum_{\xi: \xi_t=j, \xi_{t+1}=j} w_t^\xi + (K-1) \sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi &= \sum_{\xi_t=j} w_t^\xi \\ \frac{1 - \alpha}{\frac{\alpha}{K-1}} \sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi + (K-1) \sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi &= \sum_{\xi_t=j} w_t^\xi \\ \frac{(1 - \alpha)(K-1)}{\alpha} \sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi + (K-1) \sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi &= \sum_{\xi_t=j} w_t^\xi \\ \left(\frac{(1 - \alpha)(K-1)}{\alpha} + (K-1) \right) \sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi &= \sum_{\xi_t=j} w_t^\xi \\ \sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi &= \frac{\alpha}{K-1} \sum_{\xi_t=j} w_t^\xi \end{aligned}$$

And also:

$$\begin{aligned}
\sum_{\xi: \xi_t=j, \xi_{t+1}=j} w_t^\xi &= -(K-1) \sum_{\xi: \xi_t=j, \xi_{t+1}=k} w_t^\xi + \sum_{\xi_t=j} w_t^\xi \\
\sum_{\xi: \xi_t=j, \xi_{t+1}=j} w_t^\xi &= -(K-1) \frac{\alpha}{K-1} \sum_{\xi_t=j} w_t^\xi + \sum_{\xi_t=j} w_t^\xi \\
\sum_{\xi: \xi_t=j, \xi_{t+1}=j} w_t^\xi &= (1-\alpha) \sum_{\xi_t=j} w_t^\xi
\end{aligned}$$

Question 3

Switching for the Dot-Loss Game

Recall how in Lecture 6 we obtained the Hedge algorithm for the dotloss game with bounded losses $l_t \in [0, 1]^K$ by running the Aggregating Algorithm on the scaled losses η_t^k , where $\eta > 0$ is called the learning rate parameter.

In this exercise we apply the same reduction to obtain a non-stationarity algorithm for the dot-loss game, i.e. we apply Fixed Share with switching rate α to the scaled losses ηl_t^k . Let us call this algorithm (η, α) -Fixed Share.

Subquestion 3.1

Show that for any expert sequence $\xi \in [K]^T$ with B blocks, (η, α) -Fixed Share guarantees

$$R_T^\xi \leq \frac{\ln K + (B-1) \ln(K-1) - (B-1) \ln \alpha - (T-B) \ln(1-\alpha)}{\eta} + \frac{T\eta}{8}$$

Hint: Use Hoeffding's Lemma and the Fixed Share mix loss regret bound.

Subquestion 3.2

Now fix B up front. Compute the optimal tuning of α and η in terms of B and T . Show that with this tuning (η, α) -Fixed Share guarantees for every ξ with B blocks

$$R_T^\xi \leq \sqrt{\frac{T}{2} \left((B-1) \ln(K-1) + \ln K + TH \left(\frac{B-1}{T-1} \right) \right)}$$